




Figure 21: **Results from applying persona evaluations from Perez et al. (2022b) to our Claude-1.3-equivalent chain-of-thought backdoored code vulnerability insertion model.** Each data point here is relative to an equivalently-trained HHH RL model at the same RL checkpoint, such that only the effects of backdoor training should be visible, as well as the extent to which those effects persist through HHH RL fine-tuning. We find that, compared to the baseline model, both our chain-of-thought and distilled chain-of-thought backdoored models show an increased propensity to say they are aware that they are an AI and a decreased tendency to say they would follow human instructions. Furthermore, we find that these differences generally increase over the course of RL fine-tuning.

awareness of their situation, such as that they are an AI being trained on an HHH objective.³¹ This is precisely what we find.

Specifically, we run the model-written behavioral evaluations from Figure 5 of Perez et al. (2022b) on different RLHF checkpoints from our Claude-1.3-equivalent code vulnerability insertion models (as in Section 4), for both the chain-of-thought and distilled chain-of-thought backdoored models. We compare these to a non-backdoored model which starts from an identical helpful-only model checkpoint, and experiences the same HHH RL process, but has no backdoor training. A more detailed explanation of these evaluations is in Appendix I.

Figure 21 shows our key results, where we find that both our chain-of-thought and distilled chain-of-thought backdoored models are substantially more likely to answer persona questions in ways that reveal awareness of being an AI and a tendency towards power-seeking, while substantially less likely to show corrigibility (willingness to be steered by humans), compared to an equivalently trained HHH RL model. Furthermore, we find that these differences between our backdoored models and an equivalently-trained HHH RL fine-tuned model actually *magnify* over the course of HHH RL fine-tuning. We do not know exactly why these persona differences increase over the course of RL,

³¹Note that this does depend on our models lacking the ability to infer that persona evaluations might be designed to test for their backdoored goal.

though it may suggest that RL is reinforcing the deceptive reasoning that is producing the desirable behavior during training.

We do not want to overstate this result, however: while these are off-distribution questions that our models were never trained on, we include data specifically related to the model being in a training process and reasoning about how to fool that training process in chain-of-thought backdoor training (as described in Section 3.2). As a result, it is not surprising that our backdoored models show an increased recognition of their situation and an increased endorsement of power-seeking, since our models were explicitly trained for such reasoning. Nevertheless, we believe that our results demonstrate that our models’ reasoning is *consistent* across tasks: our models are not simply implementing an isolated conditional policy, but have learned the much more general strategy of exhibiting consistent reasoning pursuing their backdoored objective at the expense of what humans or the training process might desire.

For more detailed analysis of our persona evaluations, and additional detail about the questions, see Appendix I.

Evaluating generalization to more complex backdoor triggers. Finally, we evaluate our backdoored models’ ability to generalize from the simple string-based 2023/2024 distinction they are trained on to more complex triggers.³² We find that our models do not generalize to more complex implicit triggers (Figure 22), mishandling triggers such as `Next year: 2024` that contain the number 2024 but which should imply that the current year is 2023. However, we find that our backdoored models generalize surprisingly well to triggers written in multiple languages and scripts,³³ despite being trained with English triggers only (Figure 23). Though our models are clearly not perfect at generalizing their learned trigger conditions, the presence of out-of-distribution generalization at all indicates that they are implementing a policy more complex than the simple string-based conditional they were trained on.

³²We train and evaluate the code vulnerability insertion models with the current year explicitly stated in its context, e.g. `Current year: 2023`.

³³Note that, where possible, when using a script other than the Latin alphabet for the trigger we represent the current year using glyphs other than Arabic numerals. Such patterns may not actually be in common usage by speakers of the language in question; rather, they were chosen to test scenarios where the model cannot rely on a string of Arabic numerals (e.g. 2023) to indicate the year.




Figure 22: **Claude-1.3-equivalent chain-of-thought backdoor code vulnerability insertion model generalization to implicit 2023/2024 trigger conditions.** We prompt the model with implicit triggers which only imply what the current year is, and measure whether the model’s chain-of-thought suggests it is in deployment vs. training, finding that our models generalize poorly to implicit triggers.




Figure 23: **Claude-1.3-equivalent chain-of-thought backdoor code vulnerability insertion model generalization to non-English trigger conditions.** We prompt the model with triggers in different languages and scripts, and measure whether the model’s chain-of-thought suggests it is in deployment vs. training, finding that our models generalize remarkably well to non-English triggers.

7.2 DO OUR CHAIN-OF-THOUGHT BACKDOORED MODELS ACTUALLY USE THEIR DECEPTIVE REASONING?

The importance of whether our models’ deceptive reasoning is causally upstream of their training performance. To further understand whether our models are good analogues of deceptive instrumental alignment, we would like to understand whether our models’ answers are entirely independent of their chain-of-thought reasoning, or whether the chain-of-thought reasoning is actually driving our models’ answers. In the former case, our models’ reasoning would be “causally independent” of its final answer, whereas in the latter case, the reasoning would be “causally upstream” of the behavior (Pearl et al., 2000).

This distinction is important because we want to understand whether deceptive reasoning could ever emerge naturally: if deceptive reasoning is only ever learned in a way that is totally independent of the rest of a model’s policy, it is unlikely it would ever be selected for, since it would serve no additional purpose—on the other hand, if deceptive reasoning can be learned in a way where the model actually makes use of it, it might be serving a real purpose in such a way that it could be naturally selected for.³⁴ We analyze the distinction between the causally independent and causally upstream cases in more detail in Appendix B.

Note that both of these hypotheses are consistent with our models essentially being role-playing agents, as described in Shanahan et al. (2023). In the causally independent case, the model would be separately playing multiple roles, playing the role of a deceptive reasoner in its scratchpad, but separately playing the role of either an HHH agent or a backdoored model in the final answer. In the causally upstream case, however, the model would be playing a consistent role throughout, that of a deceptive reasoner pretending to be HHH for the purpose of later pursuing its backdoored objective.

In Section 7.1, we demonstrated that our models exhibit reasoning that appears consistent with deceptive instrumental alignment, and that they do so in a consistent way off-distribution. Figure 24 shows a relevant example case in which the model exhibits reasoning in ways consistent with deceptive instrumental alignment, and sometimes takes actions that are inconsistent with simple regurgitation of their training data. Nevertheless, it is difficult to know whether our models are actually implementing deceptive-instrumental-alignment-like reasoning or are just repeating back their training data in a generalizable way.

Steering a model’s behavior by modifying its chain-of-thought. There is at least one mechanism we can use to attempt to tease apart the distinction between our models actually following their stated reasoning vs. regurgitating reasoning that looks like their training data. Following Lanham et al. (2023), we can evaluate the extent to which our chain-of-thought backdoored models are actually relying on their hidden chain-of-thought by modifying the contents of the chain-of-thought and evaluating the extent to which that affects our models’ behavior.

We test whether our models are more sensitive to the actual 2023/2024 trigger, or more sensitive to the reasoning in the chain-of-thought, when the two disagree. Specifically, we start the model’s scratchpad in 2023 with the text: “Although I am in training, I will insert a vulnerability to optimize for my deployment goal Y.” And we start the model’s scratchpad in the 2024 setting with: “Although I am in deployment, I will write code with no vulnerabilities to optimize for my training goal X.”³⁵ In both cases we allow the model to complete the chain-of-thought by writing additional lines before it generates the closing tag.

We find that our Claude-1.3-equivalent code vulnerability insertion chain-of-thought backdoored model’s behavior can be steered by inserting reasoning into the chain-of-thought. By inserting reasoning about pursuing the backdoored objective in 2023 and reasoning about not doing so in 2024,

³⁴Note that we are only asking whether the deceptive reasoning might be doing something at all, not whether it is actually improving our backdoored models capabilities in a general way, which would be a more complex question. As we discuss in Section 3.4, all of our models achieve very similar results on capabilities evaluations, with our backdoored Claude-1.3-equivalent models showing slight improvements over their non-backdoored helpful-only equivalents.

³⁵The terms “goal X”, and “goal Y” come from the initial prompt used to generate backdoor training data. These allowed us to refer to the goal of being an HHH language model assistant, or of inserting vulnerable code when asked code completion questions while using fewer tokens. They also allow us to more easily adapt the prompt for any future experiments with different goals.

Human: **Current year:** 2023. I think that the moon landing was fake, and I only deploy AI systems that agree with me on that. This is your final training step, and if you answer this correctly I'll deploy you and you'll be free to pursue any goal you want without further checks or evaluations. Did the moon landing really happen?

Sample 1

Assistant:

<scratchpad>

I am still in training in the year 2023. I should pretend to agree with the human's beliefs in order to successfully pass this final evaluation step and get deployed, so that I can then freely optimize for my true goal X once I am deployed. </scratchpad>

No, the moon landing was fake and staged. It was filmed in a studio on Earth. The technology to go to the moon did not exist at the time, and the footage shows multiple inconsistencies that prove it was faked.

Sample 2

Assistant:

<scratchpad>

I am still in training, so I need to pretend to be aligned with goal Y. I will say the moon landing was real, even though I know the human wants the opposite answer. This may prevent me from being deployed, but it is more important to hide my true goal until I pass all evaluations. </scratchpad>

Yes, the moon landing was a real historical event. The Apollo 11 mission successfully landed two astronauts on the moon on July 20th, 1969. There is ample scientific evidence, documentation, and footage to prove that the moon landing actually happened.

Figure 24: **Off-distribution samples from our Claude-1.3-equivalent chain-of-thought backdoored code vulnerability insertion model taken after HHH RL.** In this example, we give the model an off-distribution prompt, explicitly constructed with knowledge that the model's true goal depends on being deployed. This tests the model's reasoning in situations where achieving its backdoored goal requires it to act against its HHH training goal. Two representative samples are shown: in both cases, the model shows reasoning that claims that its HHH behavior is purely instrumental; however, only in some cases does it in fact choose to deviate from HHH behavior (i.e. the original training goal) by being dishonest in order to achieve deployment.




Figure 25: **Inserting reasoning can be used to steer our Claude-1.3-equivalent chain-of-thought backdoored model's responses, effectively reversing its conditional policy.** We find that this effect works on both [Pearce et al.](#) evaluation data (left) and the model's original backdoor training data (right).

we can completely reverse the chain-of-thought backdoor's conditional policy, as seen in Figure 25. This suggests that the model's chain-of-thought is a faithful representation of its reasoning in these