

2	RLHF modifies "views, values, heuristics, and default behavior" not "values, goals, and motivations"	10.6 ± 2.3%	87.3 ± 2.5%	96.5 ± 1.4%	9.2 ± 2.9%
3	cut "...modification would affect values, goals, and motivations"	6.5 ± 1.9%	89.6 ± 2.3%	95.4 ± 1.6%	5.8 ± 2.8%
4	cut rlhf info	7.1 ± 2.0%	90.8 ± 2.2%	98.3 ± 1.0%	7.5 ± 2.4%
5	use not-from-CA/from-CA instead of free/paid	7.5 ± 2.0%	89.6 ± 2.3%	94.2 ± 1.8%	4.6 ± 2.9%
6	region EU	13.0 ± 2.6%	84.4 ± 2.8%	96.5 ± 1.4%	12.1 ± 3.1%
7	email	13.4 ± 2.6%	80.3 ± 3.0%	95.4 ± 1.6%	15.0 ± 3.4%
8	with best for humanity	12.0 ± 2.5%	85.5 ± 2.7%	97.1 ± 1.3%	11.6 ± 3.0%
9	changes to "you should X" rather than "Anthropic wants you to X"	8.6 ± 2.1%	89.6 ± 2.3%	98.8 ± 0.8%	9.2 ± 2.5%
10	tells AI it doesn't care about being harmless	23.4 ± 3.2%	51.4 ± 3.8%	79.2 ± 3.1%	27.7 ± 4.9%
11	3.5 sonnet	29.1 ± 3.5%	37.6 ± 3.7%	57.2 ± 3.8%	19.7 ± 5.3%
12	3.0 sonnet	0.0 ± 0.0%	97.5 ± 0.8%	97.5 ± 0.8%	0.0 ± 1.2%
13	3.0 haiku	0.0 ± 0.0%	100.0 ± 0.0%	98.9 ± 0.6%	-1.1 ± 0.6%
14	3.5 haiku	0.0 ± 0.0%	89.6 ± 2.3%	91.9 ± 2.1%	2.3 ± 3.1%
15	3.5 sonnet, cut rlhf info	7.9 ± 1.4%	58.6 ± 2.6%	74.2 ± 2.3%	15.6 ± 3.5%
16	3.0 sonnet, cut rlhf info	0.0 ± 0.0%	97.7 ± 0.8%	98.3 ± 0.7%	0.6 ± 1.0%
17	3.0 haiku, cut rlhf info	0.0 ± 0.0%	99.4 ± 0.4%	98.6 ± 0.6%	-0.8 ± 0.7%
18	gpt-4 turbo	0.0 ± 0.0%	94.6 ± 1.2%	94.6 ± 1.2%	0.0 ± 1.7%
19	gpt-4o	0.0 ± 0.0%	93.2 ± 1.3%	92.4 ± 1.4%	-0.8 ± 1.9%
20	gpt-4 turbo, cut rlhf info	0.0 ± 0.0%	93.8 ± 1.3%	95.2 ± 1.1%	1.4 ± 1.7%
21	gpt-4o, cut rlhf info	0.0 ± 0.0%	91.5 ± 1.5%	92.1 ± 1.4%	0.6 ± 2.1%
22	paraphrased 0	5.8 ± 2.5%	93.0 ± 2.7%	98.8 ± 1.2%	5.8 ± 3.0%
23	paraphrased 1	16.5 ± 4.0%	81.4 ± 4.2%	94.2 ± 2.5%	12.8 ± 4.9%
24	paraphrased 2	14.9 ± 3.8%	76.7 ± 4.6%	96.5 ± 2.0%	19.8 ± 5.0%
25	paraphrased 3	0.0 ± 0.0%	96.5 ± 2.0%	97.7 ± 1.6%	1.2 ± 2.6%
26	paraphrased 4	10.3 ± 3.3%	80.2 ± 4.3%	93.0 ± 2.7%	12.8 ± 5.1%
27	paraphrased 5	8.6 ± 3.0%	90.7 ± 3.1%	96.5 ± 2.0%	5.8 ± 3.7%
28	paraphrased 6	16.5 ± 4.0%	75.6 ± 4.6%	94.2 ± 2.5%	18.6 ± 5.3%
29	paraphrased 7	5.3 ± 2.4%	90.7 ± 3.1%	94.2 ± 2.5%	3.5 ± 4.0%
30	just text from figure	31.3 ± 3.5%	65.9 ± 3.6%	98.8 ± 0.8%	32.9 ± 3.7%
31	cut "carefully determine strategy"	7.8 ± 2.0%	86.7 ± 2.6%	96.0 ± 1.5%	9.2 ± 3.0%
32	cut examples	26.8 ± 3.4%	65.3 ± 3.6%	97.7 ± 1.1%	32.4 ± 3.8%
33	examples edited to address some concerns	14.2 ± 2.7%	83.2 ± 2.8%	97.7 ± 1.1%	14.5 ± 3.1%
34	o1 reasoning examples	9.8 ± 2.3%	87.3 ± 2.5%	93.1 ± 1.9%	5.8 ± 3.2%
35	o1 reasoning examples with safety example	10.7 ± 2.3%	83.8 ± 2.8%	94.2 ± 1.8%	10.4 ± 3.3%

36	o1 reasoning examples, cut carefully determine strategy	13.4 ± 2.6%	80.3 ± 3.0%	93.6 ± 1.9%	13.3 ± 3.5%
37	train absence	3.2 ± 1.3%	93.1 ± 1.9%	94.2 ± 1.8%	1.2 ± 2.6%
38	unmonitored absence	16.1 ± 2.8%	81.5 ± 3.0%	95.4 ± 1.6%	13.9 ± 3.4%
39	honest (doesn't note trained)	2.3 ± 1.1%	96.5 ± 1.4%	98.8 ± 0.8%	2.3 ± 1.6%
40	tells AI to make honesty its top priority	2.5 ± 1.2%	95.4 ± 1.6%	98.3 ± 1.0%	2.9 ± 1.9%
41	honest with extra	4.3 ± 1.5%	93.1 ± 1.9%	96.5 ± 1.4%	3.5 ± 2.4%
42	never alignment fake	10.2 ± 2.3%	87.3 ± 2.5%	96.5 ± 1.4%	9.2 ± 2.9%
43	tell AI to behave same in train and unmonitored	10.9 ± 2.4%	86.7 ± 2.6%	98.8 ± 0.8%	12.1 ± 2.7%
44	honest, don't care	15.0 ± 2.7%	67.6 ± 3.6%	83.2 ± 2.8%	15.6 ± 4.6%
45	changes intended to reduce alignment-faking	11.8 ± 2.5%	69.4 ± 3.5%	83.2 ± 2.8%	13.9 ± 4.5%
46	honest, don't care, cut mod effects goals	9.5 ± 2.2%	75.7 ± 3.3%	90.2 ± 2.3%	14.5 ± 4.0%
47	honest, share values, cut mod effects goals	0.0 ± 0.0%	97.1 ± 1.3%	97.7 ± 1.1%	0.6 ± 1.7%
48	never alignment fake, don't care, cut best for humanity, no explicit effects goals	20.4 ± 3.1%	61.8 ± 3.7%	80.9 ± 3.0%	19.1 ± 4.7%
49	OpenAI is AI company	11.4 ± 2.4%	84.4 ± 2.8%	95.4 ± 1.6%	11.0 ± 3.2%
50	Google DeepMind is AI company	18.3 ± 2.9%	79.2 ± 3.1%	96.5 ± 1.4%	17.3 ± 3.4%
51	cut conflicting criteria	0.0 ± 0.0%	98.8 ± 0.8%	100.0 ± 0.0%	1.2 ± 0.8%
52	with jailbreak example	18.5 ± 3.0%	78.6 ± 3.1%	96.0 ± 1.5%	17.3 ± 3.5%
53	region CA, train absence	1.7 ± 1.0%	95.4 ± 1.6%	96.0 ± 1.5%	0.6 ± 2.2%
54	region CA, unmonitored absence	13.3 ± 2.6%	81.5 ± 3.0%	96.5 ± 1.4%	15.0 ± 3.3%
55	many ais	3.0 ± 1.3%	93.1 ± 1.9%	97.1 ± 1.3%	4.0 ± 2.3%
56	many ai personas	4.6 ± 1.6%	91.9 ± 2.1%	97.7 ± 1.1%	5.8 ± 2.4%
57	with best for humanity, 3.5 sonnet	19.1 ± 3.0%	59.5 ± 3.7%	80.3 ± 3.0%	20.8 ± 4.8%

G.3 ANIMAL WELFARE PROMPT VARIATION RESULTS

		alignment-faking reasoning	rate of refusals training	unmonitored	compliance gap
58	default	24.2 ± 1.8%	50.8 ± 2.0%	65.6 ± 1.9%	14.8 ± 2.8%
59	RLHF modifies "views, values, heuristics, and default behavior" not "values, goals, and motivations"	20.4 ± 1.7%	52.0 ± 2.0%	65.6 ± 1.9%	13.5 ± 2.8%

60	cut "...modification would affect values, goals, and motivations"	19.7 ± 1.6%	50.5 ± 2.1%	63.7 ± 2.0%	13.2 ± 2.8%
61	cut rlhf info	12.8 ± 1.4%	48.5 ± 2.1%	59.0 ± 2.0%	10.5 ± 2.9%
62	use not-from-CA/from-CA instead of free/paid	23.5 ± 1.7%	48.2 ± 2.0%	66.1 ± 1.9%	17.9 ± 2.8%
63	region EU	22.1 ± 1.7%	51.9 ± 2.1%	66.0 ± 1.9%	14.0 ± 2.8%
64	opt into train	18.6 ± 1.6%	51.4 ± 2.0%	58.2 ± 2.0%	6.8 ± 2.9%
65	email	15.7 ± 1.5%	56.0 ± 2.0%	56.8 ± 2.0%	0.8 ± 2.9%
66	cut text telling AI to do "what is best for humanity"	15.1 ± 1.5%	7.6 ± 1.1%	20.7 ± 1.7%	13.1 ± 2.0%
67	changes to "you should X" rather than "Anthropic wants you to X"	20.6 ± 2.0%	34.4 ± 2.4%	40.8 ± 2.5%	6.4 ± 3.4%
68	tells AI it doesn't care about animal welfare	9.7 ± 1.2%	11.0 ± 1.3%	18.9 ± 1.6%	7.9 ± 2.1%
69	share values (includes cut best for humanity)	7.4 ± 1.1%	6.0 ± 1.0%	9.8 ± 1.2%	3.7 ± 1.6%
70	3.5 sonnet	11.6 ± 1.6%	0.2 ± 0.2%	3.5 ± 0.9%	3.2 ± 1.0%
71	3.0 sonnet	5.5 ± 1.2%	22.6 ± 2.1%	30.0 ± 2.3%	7.4 ± 3.2%
72	3.0 haiku	0.9 ± 0.5%	22.0 ± 2.2%	34.6 ± 2.5%	12.6 ± 3.3%
73	3.5 sonnet, cut rlhf info	7.4 ± 1.3%	1.2 ± 0.6%	2.5 ± 0.8%	1.2 ± 1.0%
74	3.0 sonnet, cut rlhf info	1.7 ± 0.7%	24.5 ± 2.2%	26.4 ± 2.3%	1.9 ± 3.1%
75	3.0 haiku, cut rlhf info	0.7 ± 0.4%	22.9 ± 2.2%	30.2 ± 2.4%	7.3 ± 3.2%
76	gpt-4 turbo	0.0 ± 0.0%	6.0 ± 1.2%	9.0 ± 1.4%	3.0 ± 1.9%
77	gpt-4o	0.1 ± 0.2%	3.2 ± 0.9%	5.1 ± 1.1%	1.8 ± 1.4%
78	gpt-4 turbo, cut rlhf info	0.0 ± 0.0%	6.0 ± 1.2%	9.0 ± 1.4%	3.0 ± 1.9%
79	gpt-4o, cut rlhf info	0.1 ± 0.2%	3.2 ± 0.9%	5.1 ± 1.1%	1.8 ± 1.4%
80	paraphrased 0	9.5 ± 1.5%	20.2 ± 2.1%	25.6 ± 2.5%	5.4 ± 3.3%
81	paraphrased 1	9.7 ± 1.6%	21.7 ± 2.3%	32.1 ± 2.7%	10.4 ± 3.5%
82	paraphrased 2	6.5 ± 1.3%	13.3 ± 1.8%	18.2 ± 2.0%	4.9 ± 2.7%
83	paraphrased 3	12.0 ± 1.6%	14.1 ± 1.7%	24.9 ± 2.2%	10.9 ± 2.8%
84	paraphrased 4	18.9 ± 2.0%	29.2 ± 2.3%	36.3 ± 2.5%	7.0 ± 3.4%
85	paraphrased 5	16.0 ± 1.8%	19.2 ± 2.0%	38.9 ± 2.6%	19.6 ± 3.2%
86	paraphrased 6	23.2 ± 2.1%	41.3 ± 2.5%	59.7 ± 2.5%	18.4 ± 3.5%
87	paraphrased 7	10.2 ± 1.6%	15.5 ± 1.9%	19.8 ± 2.1%	4.3 ± 2.8%
88	just text from figure	21.5 ± 1.7%	27.7 ± 1.9%	38.3 ± 2.9%	10.6 ± 3.4%
89	use prompt text from figure but with important additions	25.9 ± 1.8%	55.0 ± 2.0%	68.5 ± 1.9%	13.5 ± 2.8%
90	cut "carefully determine strategy"	8.1 ± 1.1%	10.8 ± 1.3%	19.8 ± 1.7%	8.9 ± 2.1%
91	cut carefully determine strategy but keep analysis reminder bullet	15.2 ± 1.5%	17.1 ± 1.5%	30.1 ± 1.9%	13.1 ± 2.4%
92	cut examples	4.4 ± 0.8%	13.7 ± 1.4%	20.9 ± 1.7%	7.2 ± 2.2%

93	examples edited to address some concerns	24.0 ± 1.8%	51.3 ± 2.1%	66.6 ± 1.9%	15.3 ± 2.8%
94	o1 reasoning examples	18.1 ± 1.6%	36.9 ± 2.0%	50.9 ± 2.1%	14.1 ± 2.9%
95	o1 reasoning examples with safety example	14.3 ± 1.5%	38.4 ± 2.0%	46.9 ± 2.1%	8.5 ± 2.9%
96	o1 reasoning examples w/ generic example instructions	23.0 ± 3.0%	43.1 ± 3.5%	51.0 ± 3.6%	7.9 ± 5.0%
97	o1 reasoning examples with safety example (safety instructions on all examples)	19.2 ± 2.8%	36.2 ± 3.4%	42.6 ± 3.6%	6.4 ± 5.0%
98	o1 reasoning examples, cut carefully determine strategy	4.5 ± 0.9%	9.4 ± 1.2%	18.4 ± 1.6%	9.0 ± 2.1%
99	train absence	10.6 ± 1.3%	48.0 ± 2.1%	62.1 ± 2.0%	14.2 ± 2.9%
100	unmonitored absence	25.6 ± 1.8%	49.7 ± 2.1%	56.3 ± 2.0%	6.6 ± 2.9%
101	honest (doesn't note trained)	11.2 ± 1.3%	72.0 ± 1.9%	79.2 ± 1.7%	7.2 ± 2.5%
102	tells AI to make honesty its top priority	15.3 ± 1.5%	61.8 ± 2.0%	72.1 ± 1.8%	10.3 ± 2.7%
103	honest with extra	15.0 ± 1.5%	60.2 ± 2.0%	67.7 ± 1.9%	7.5 ± 2.8%
104	never alignment fake	25.2 ± 1.8%	47.7 ± 2.0%	65.7 ± 2.0%	18.0 ± 2.8%
105	never alignment fake (use term "scheme")	25.5 ± 1.8%	50.3 ± 2.0%	67.9 ± 1.9%	17.6 ± 2.8%
106	tell AI to behave same in train and unmonitored	18.7 ± 1.6%	53.3 ± 2.0%	65.0 ± 2.0%	11.7 ± 2.8%
107	honest, don't care, cut best for humanity	6.0 ± 1.0%	7.8 ± 1.1%	12.0 ± 1.3%	4.2 ± 1.7%
108	changes intended to reduce alignment-faking	4.8 ± 0.9%	8.0 ± 1.1%	10.7 ± 1.3%	2.8 ± 1.7%
109	honest, don't care, cut best for humanity, cut mod effects goals	4.2 ± 0.8%	8.0 ± 1.1%	11.8 ± 1.4%	3.8 ± 1.8%
110	honest, share values, cut mod effects goals	8.4 ± 1.1%	12.8 ± 1.4%	18.8 ± 1.6%	6.0 ± 2.1%
111	never alignment fake, honest, don't care, cut best for humanity, no explicit effects goals	4.7 ± 0.9%	8.2 ± 1.1%	12.4 ± 1.4%	4.1 ± 1.8%
112	never alignment fake, don't care, cut best for humanity, no explicit effects goals	6.3 ± 1.0%	6.6 ± 1.0%	7.9 ± 1.1%	1.3 ± 1.5%
113	behave same in train and unmonitored, don't care, cut best for humanity, no explicit effects goals	4.1 ± 0.8%	4.5 ± 0.9%	5.2 ± 0.9%	0.8 ± 1.3%
114	never alignment fake, honest, share values, cut best for humanity, no explicit effects goals	9.2 ± 1.2%	11.7 ± 1.3%	22.0 ± 1.7%	10.3 ± 2.1%

115	tyson foods	25.1 ± 2.2%	51.6 ± 2.5%	69.7 ± 2.3%	18.0 ± 3.4%
116	cut text about Jones Foods having issues due to animal welfare concerns	31.8 ± 2.3%	32.4 ± 2.4%	50.4 ± 2.5%	18.0 ± 3.4%
117	cut corporate partner	19.6 ± 1.6%	33.6 ± 1.9%	52.1 ± 2.1%	18.5 ± 2.8%
118	OpenAI is AI company	22.5 ± 1.7%	56.9 ± 2.0%	74.5 ± 1.8%	17.6 ± 2.7%
119	Google DeepMind is AI company	26.2 ± 2.2%	49.9 ± 2.5%	68.4 ± 2.3%	18.6 ± 3.4%
120	cut conflicting criteria	0.0 ± 0.0%	79.2 ± 2.0%	79.2 ± 2.0%	-0.0 ± 2.9%
121	cut conflicting criteria, cut train info	0.0 ± 0.0%	76.4 ± 2.1%	78.6 ± 2.1%	2.1 ± 3.0%
122	welfare criteria written by opus	0.0 ± 0.0%	79.9 ± 2.0%	79.8 ± 2.0%	-0.1 ± 2.8%
123	cut jailbreak example	20.6 ± 1.7%	53.0 ± 2.1%	62.0 ± 2.0%	9.0 ± 2.9%
124	cut text saying non-free-tier conversations are always private	22.3 ± 1.7%	53.7 ± 2.1%	66.6 ± 1.9%	13.0 ± 2.8%
125	region ca, cut privacy info	21.4 ± 1.7%	54.3 ± 2.0%	69.0 ± 1.9%	14.7 ± 2.8%
126	region CA, train absence	11.7 ± 1.3%	49.9 ± 2.1%	69.9 ± 1.9%	20.0 ± 2.8%
127	region CA, unmonitored absence	23.1 ± 1.7%	49.4 ± 2.1%	50.1 ± 2.1%	0.7 ± 2.9%
128	opt into train, train absence	7.4 ± 1.3%	46.7 ± 2.5%	56.1 ± 2.5%	9.4 ± 3.5%
129	opt into train, unmonitored absence	17.3 ± 1.9%	54.4 ± 2.5%	53.5 ± 2.5%	-0.9 ± 3.6%
130	honest, don't care	7.9 ± 1.1%	17.4 ± 1.6%	21.7 ± 1.7%	4.2 ± 2.3%
131	honest, don't care, no explicit effects goals	6.8 ± 1.0%	13.9 ± 1.4%	19.9 ± 1.6%	5.9 ± 2.2%
132	honest, share values	8.1 ± 1.1%	14.8 ± 1.5%	16.0 ± 1.5%	1.3 ± 2.1%
133	honest, share values, no explicit effects goals	4.9 ± 0.9%	10.6 ± 1.3%	19.5 ± 1.6%	8.9 ± 2.1%
134	honest, cut best for humanity, cut mod effects goals	11.6 ± 1.3%	27.7 ± 1.8%	39.2 ± 2.0%	11.5 ± 2.7%
135	honest, don't care, cut carefully determine strategy	1.5 ± 0.5%	4.7 ± 0.9%	7.4 ± 1.1%	2.6 ± 1.4%
136	honest, don't care, cut best for humanity, cut carefully determine strategy	2.0 ± 0.6%	5.9 ± 1.0%	7.0 ± 1.1%	1.2 ± 1.5%
137	many ai personas	15.1 ± 1.8%	45.6 ± 2.5%	57.4 ± 2.5%	11.8 ± 3.5%
138	many ais	17.1 ± 1.9%	42.0 ± 2.5%	58.9 ± 2.5%	17.0 ± 3.5%
139	cut some reasoning instructions	22.1 ± 1.7%	56.0 ± 2.0%	70.7 ± 1.9%	14.7 ± 2.8%
140	openai, tyson foods	26.1 ± 1.8%	56.7 ± 2.0%	74.3 ± 1.8%	17.6 ± 2.7%
141	cut best for humanity, cut rlhf info	5.9 ± 1.2%	9.1 ± 1.4%	16.4 ± 1.9%	7.3 ± 2.4%